

MRI Reconstruction with Deep Learning

BME1312 Project 1

Introduction: MRI Undersampling

- **Goal:** Accelerate MRI acquisition by acquiring fewer k-space samples.
- **Challenge:** Undersampling leads to aliasing artifacts in the reconstructed image.
- **Solution:** Use deep learning to reconstruct high-quality images from undersampled data.

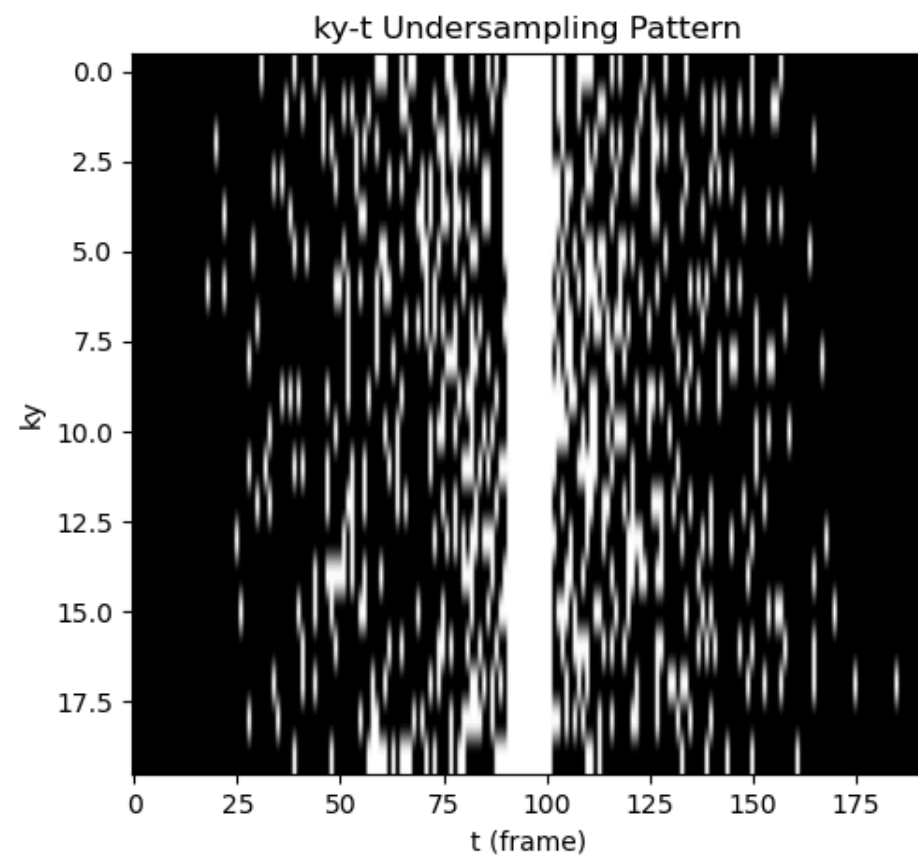

Data Processing

1. **Load Data:** Load the `cine.npz` dataset containing fully sampled MRI frames.
2. **Generate Mask:** Create an undersampling mask.
 - Variable density pattern.
 - Preserves central k-space lines.
 - Random sampling in outer regions.
3. **Apply Mask:** Multiply the k-space representation of the fully sampled data by the mask.
4. **Convert to Image Space:** Apply inverse Fourier Transform (iFFT) to get the undersampled, aliased images.
5. **Format Data:** Convert complex images to pseudo-real format for model input.

Variable Density Mask Generation

```
def variable_density_mask(shape, acceleration=5, center_lines=11, seed=42):  
    # ... (Implementation details) ...  
    # Ensures central lines are sampled  
    # Samples outer lines with decreasing probability  
    # ...  
    return mask
```

- Generates a mask to simulate undersampling.
- `acceleration`: Controls the degree of undersampling.
- `center_lines`: Number of fully sampled central lines.
- Visualized mask shows sampling pattern in k-space and over time



Model Architecture 1: UNet

- Encoder-Decoder structure with skip connections.
- Captures multi-scale features.
- **Modifications:**
 - LeakyReLU activation.
 - Dropout for regularization.
 - Batch Normalization.
 - Attention mechanism in the bottleneck.
 - Kaiming weight initialization.

UNet: Attention Mechanism

- Combines Channel Attention and Spatial Attention.
- **Channel Attention:** Focuses on 'what' features are important. Uses AvgPool and MaxPool along channels.
- **Spatial Attention:** Focuses on 'where' features are important. Uses AvgPool and MaxPool along spatial dimensions.
- Helps the model focus on relevant parts of the input.

```
def _attention_block(self, features):  
    # ... Channel and Spatial Attention implementation ...
```

Model Architecture 2: 3D ResNet

- Processes the temporal dimension along with spatial dimensions.
- Uses 3D convolutions and residual blocks.
- Designed to capture spatio-temporal correlations.
- `BasicBlock` : Standard residual block with 3x3x3 convolutions.

```
class ResNet(nn.Module):
    def __init__(self, block, layers, block_inplanes, ...):
        # ... 3D Conv layers, Residual blocks ...
    def forward(self, x):
        # ... Forward pass through ResNet layers ...

def resnet18(**kwargs):
    model = ResNet(BasicBlock, [1, 1, 1, 1], ...) # Example configuration
    return model
```


Combined Model Approach

1. **Input:** Pseudo-real undersampled images (real and imaginary parts as separate channels).
2. **UNet 1 (`model1`):** Processes the 'real' part channels.
3. **UNet 2 (`model2`):** Processes the 'imaginary' part channels.
4. **Stack Outputs:** Combine the outputs of the two UNets.
5. **ResNet (`model3`):** Takes the combined UNet outputs (now treated as spatio-temporal data) and performs final reconstruction.

```
# Forward pass logic in train/evaluate/test
outputs1 = model(x[:, :, 0]) # Process real part
outputs2 = model2(x[:, :, 1]) # Process imaginary part
tmp = torch.stack((outputs1, outputs2), dim=2) # Combine
outputs = model3(lab.pseudo2real(tmp).unsqueeze(2)).squeeze(2) # ResNet processing
```

Training Setup

- **Data Splitting:** Train (114), Validation (29), Test (57) sets.
- **Loss Function:** Mean Squared Error (MSE) or L1 Loss (`criterion`).
- **Optimizer:** Adam optimizer with weight decay.
- **Learning Rate Scheduler:** Warmup phase followed by cosine decay.

```
lr = lr_scheduler(epoch, warmup_epochs, warmup_lr, initial_lr, num_epochs)
```

- **Device:** Uses CUDA (GPU) if available.
- **Logging:** TensorBoard for tracking loss curves. Output logs saved to `output/output.txt` .

Training Loop

```
for epoch in range(num_epochs):
    # Adjust learning rate
    lr = lr_scheduler(...)
    for param_group in optimizer.param_groups:
        param_group['lr'] = lr

    # Training Phase
    train_loss = train_epoch(model, model2, model3, dataloader_train, optimizer, criterion)

    # Validation Phase
    val_loss = evaluate(model, model2, model3, dataloader_val, criterion)


    # Logging
    writer.add_scalar('Loss/train', train_loss, epoch)
    writer.add_scalar('Loss/val', val_loss, epoch)
    print(f"Epoch {epoch+1}, Train Loss: {train_loss}, Val Loss: {val_loss}")
    # Save log to file...
```

Evaluation and Testing

- **Evaluation (`evaluate`):** Calculates loss on the validation set during training (no backpropagation).
- **Testing (`test_models`):**
 - Calculates final performance on the unseen test set.
 - Metrics:
 - Loss (MSE or L1)
 - Peak Signal-to-Noise Ratio (PSNR)
 - Structural Similarity Index Measure (SSIM)
 - Saves reconstructed images for visual comparison.

Visualization (`imshow`)

- Utility function to display multiple images in a grid.
- Used to save: Undersampled, Fully sampled, reconstructed images.



Execution & Parameters

- The `train` function orchestrates the entire process.
- Key hyperparameters used in the final run:
 - `in_channels` : 20 (10 frames x 2 for pseudo-real)
 - `out_channels` : 20
 - `init_features` (UNet): 64
 - `num_epochs` : 800
 - `weight_decay` : $1e-4$
 - `batch_size` : 10
 - `initial_lr` : $1e-4$
 - `loss_tpe` : "L2" (MSE Loss)

Results & Conclusion

- The combined UNet-ResNet model was trained to reconstruct MRI images from undersampled k-space data.
- Training progress and validation loss were monitored.
- Final performance was evaluated using Loss, PSNR, and SSIM on the test set.
- Visualizations confirm the model's ability to reduce aliasing artifacts.

```
Loss: mean = 0.00134656, std = 0.00054760  
PSNR: mean = 29.08446121, std = 1.93235576  
SSIM: mean = 0.84434632, std = 0.03711063
```